

Assignment - 02.

Ans. 1. Address Translation

- CPU uses Virtual Address (VA)
- MMU splits VA → page number + offset.
- TLB lookup : if hit → directly gives Physical frame → Physical Address (PA)
- If miss → page table walk → update TLB → get PA.
- Each process has its own page table to ensure isolation and protection.

Ans. 2. Fragmentation

- Internal : unused space inside allocated block (e.g. demand up pages)
- External : scattered free blocks prevent large requests.
- Fine (beyond compaction) : Paging, segmentation with paging, buddy allocation, slab, objects for small data, huge page, pure char page allocation

Ans. 3. Paging mode and trade off.

- Memory divided into fixed-size pages / frames.
- Page table + TLB handle mapping
- Pros : no external fragmentation, flexible allocation.
- Cons : page table memory overhead, TLB misses and delay, small interval fragmentation inside page.

Ans. 4. OS - hardware interaction.

- Hardware : TLB, MMU, PTEs with R/W/X bits, ASID, page fault mechanism.
- OS : maintains page table, handles page faults,

manages TLB shutdown.
Example: invalid access \rightarrow trap \rightarrow OS loads page \rightarrow update PTE / TLB \rightarrow process.

Ques 5: Virtual address width = 16 bits = total virtual
- all address space = 2^{16} bytes.
 $2^{16} = 65,536$ bytes.

$$\rightarrow \text{Page size} = 1\text{KB} = 1024 \text{ bytes} = 2^{10} \text{ bytes}.$$

$$\rightarrow \text{Number of pages} = \frac{2^{16}}{2^{10}} = 2^6 = 64 \text{ pages}.$$

$$\rightarrow \text{Page Table size (each PTE = 2 bytes)}: 64 \times 2 \text{ bytes} \\ = 128 \text{ bytes}.$$

Answer - 64 pages, page table = 128 bytes.

Ques 6: (stepwise)

① Allocate P1 (212 KB) from the 1000 KB block.

$$\text{Remaining free} = 1000 - 212 = 788 \text{ KB}.$$

$$P1: 212, \text{ free} \rightarrow 788 \text{ KB}.$$

② Allocate P2 (417 KB) from the first free block
(788 KB)

$$\text{Remaining free} = 788 - 417 = 371 \text{ KB}.$$

$$[P1: 212] [P2: 417] [\text{free: } 371].$$

③ Allocate P3 (112 KB) from the first free
block (371 KB)

$$\text{Remaining free} = 371 - 112 \\ = 259 \text{ KB}.$$

$$[P1: 212] [P2: 417] [P3: 112] [\text{free: } 259].$$

④ Attempt P4 (426 KB): only free block is
259 cannot allocate.

\rightarrow After P1, P2, P3 allocation $\rightarrow 259 \text{ KB}$ unused;
P4 cannot be allocated.

Ans. 7

Reference string : 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0,
frame = 3.
3, 2 (13 egs)

→ FIFO : 10 faults.

→ LRU : 9 faults.

→ Optimal : 7 faults. (minimum)

Observation → Optimal is best. LRU is close to optimal and better than FIFO. FIFO may suffer Belady's anomaly, but LRU & optimal (stack algorithms) do not.

Ans. 8. (a) Overhead.

→ Dirty pages : $0.3 \times 1000 = 300$

→ Disk writes : $300 \times 10 \text{ ms} = 3000 \text{ ms} = 31 \text{ s}$.

→ Memory writes : 0.1 ms (negligible)
Latency overhead ≈ 3 seconds.

(b) Optimizations.

→ Prefer clean pages as victim.

→ Asynchronous writes behind

→ Batch / coalesce writes to reduce latency.

→ Compression before writeback.

→ Use SSD / NVMe for faster writes.

→ Use ~~SSD~~ / NVMe for faster writes.

→ Use SSD / NVMe for faster writes.

→ Reduce Replacements.

Ans. 9. Internally cache memory.

a) Use working set model: keep mission critical pages pinned; replace pages from infotainment first.

b) Strategy: use memory pool for real-time tasks. Use dynamic pool for others. adopt page and prefetching for efficiency.